

julia> Lecture Six: Parallelization and UW Resources

Course: Computational Bootcamp

Name: John Higgins

Date: July 3, 2024

```
julia> Today's goals
```

1. Understand what parallel computing is and the basics of how it works
2. Understand how to utilize it to assist with research problems

julia> Parallelization

- > Central Processing Unit: Executes instructions of a computer program
- > Modern CPUs have multiple processor cores (“processes”)
 - > Different processes can execute different tasks at same time (multi-processing)
 - > Each core has its own memory
- > Modern cores/CPU execute several tasks at the same time (multi-threading)
 - > Each task is called a thread
 - > Threads will share memory
- > Intuition:
 - > Multi-threading: multiple people (cores) baking pizza (executing code) in the *same* kitchen (CPU)
 - > Multi-processing: multiple people (cores) baking pizza (executing code) in *separate* kitchens (CPU)

julia> Multi-processing or multi-threading?

- > When is multi-threading better than multi-processing?
- > Multi-threading:
 - > All workers share same memory
 - > Faster set-up and faster to share memory across workers
 - > Need to be careful when workers edit same memory (“data race”)
 - > All workers need to be on same computer/server
- > Multi-processing:
 - > All workers have their own memory (kitchen, ingredients)
 - > Slower to set up and share memory
 - > Need to be careful about what information each worker needs
 - > Workers can exist across computers/servers

julia> Multi-processing or multi-threading?

> General rule of thumb:

- > If each task is relatively quick/small, multi-threading may be more efficient
- > If each task is quite intensive, multi-processing may be more efficient

> To belabor the pizza analogy:

- > It is cheaper to only run one kitchen (less equipment, less duplication), but can be inefficient (limited cooking space, limited ovens, more congestion)
- > It is more expensive to operate multiple kitchens, but each kitchen operates independently and is less congested

julia> What does this mean for you?

- > Multi-processing and multi-threading allow you to complete tasks in parallel and speed up your code if done well
- > So far, you are not taking full advantage of computation resources available to you
- > Your laptop likely has several cores that can have two threads of execution
- > UW-Madison has a lot of computing power available for you to use

julia> Social Science Computing Cooperative

> <https://sscc.wisc.edu/sscc/pubs/grc/resources.html>

1. Linstat: 4 interactive Linux servers (

> Each server has 36 or 48 cores, 500GB of memory

> Good for checking that your code actually works on SSCC's servers before submitting to Slurm

2. Slurm: 40 Servers with a total of 5,000 cores

>

> Most servers have 128 cores and at least 250GB of memory (1,024GB)

> You can let your code run for up to 10 days

> Be mindful that other people need to use Linstat and Slurm, too

> SSCC has other resources for jobs that need lots of memory (Winstat BIG) and jobs that use confidential data (SILO)

> If you want to use the power of Linstat in an interactive environment like a Windows/Mac, you can use the SSCC's new Open OnDemand service:

<https://kb.wisc.edu/sscc/page.php?id=131429>

julia> Center for High Throughput Computing

- > <https://chtc.cs.wisc.edu/>
- > High Performance Computing (HPC): Multi-processing tasks
- > High Throughput Computing Cluster (HTC): Multi-threading tasks
- > I am not so familiar with these resources

julia> Example: Solving capital savings

> Value function iteration to solve for V :

$$V_{n+1}(z, k) = \max_{c, k'} \{ \log(c) + \beta E_{z'} [V_n(z', k') \mid z] \}$$

$$c + k' = zk^\alpha + (1 - \delta)k$$

- > Solving on some capital grid $\{k_1, \dots, k_n\}$
- > To use parallelization, we can ask different cores/workers to solve for $\{k'(z, k), c(z'k), V_{n+1}(z'k)\}$ at different initial k
- > Then, we gather the information together and complete the next iteration

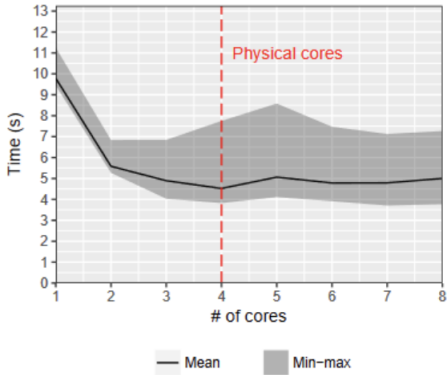
julia> Example: Simulating capital savings

- > Now we have solved for V
- > Suppose we want to simulate lots of histories of c and k
- > We generate a bunch of random paths for productivity $\{z_t\}$
- > Simulate paths of c and k
- > Instead of doing this all on one core, we can simulate many different paths at the same time using parallel computing

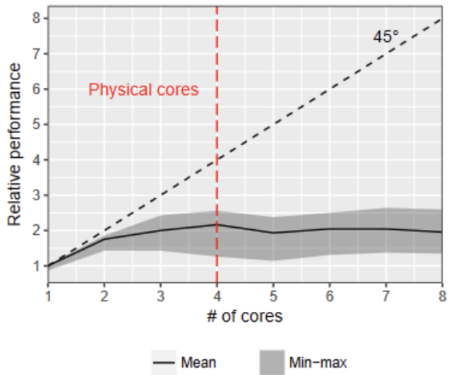
julia> No free lunch

- > Parallelization is not a free lunch because of over-head costs
- > Having many cores running uses more memory
- > Sharing memory across cores takes time
- > If you are doing lots of calculations that take a long time to run, parallelization can help
- > There are also diminishing returns that can kick in quickly for smaller problems

julia> Simple savings model with many cores



(a) Computing time (s)



(b) Performance gains.

Figure 8: Results in Julia with different number of processors. Number of experiments: 20.

https://www.sas.upenn.edu/~jesusfv/Guide_Parallel.pdf

julia> Accessing Linstat and Slurm

- > The easiest way to log in to Linstat/Slurm is to first log on to Winstat
- > You should have made a Winstat account/password at first-year orientation
- > <https://kb.wisc.edu/sscc/using-winstat>
- > You can log into Winstat with Citrix Workspace or at the following website. Note: Web log-in cannot access local files.
(<https://winstat.ssc.wisc.edu/vpn/index.html>)
- > If you need to reset your password:
https://www.ssc.wisc.edu/sscc_jsp/password/reset.jsp
- > Note: you can also log into Linstat directly from your computer
 - > <https://sscc.wisc.edu/sscc/pubs/grc/linux.html>
 - > On Windows, need to install X-Win32
 - > On Mac, can connect directly using command line
 - > On either OS, you will need to use the SSCC VPN
(<https://kb.wisc.edu/sscc/connecting-to-the-sscc-network-via-vpn>)

julia> Linstat Cheat Sheet

- > Open interactive Julia REPL:

```
julia
```

- > Run a Julia script file (titled “script.jl”) and save output to text file (titled “output.txt”):

```
julia script.jl > output.txt &
```

- > Change working directory to a specified folder (given by path “FolderPath”
- e.g. FolderPath = “/project_folder/code/model/”):

```
cd FolderPath
```

julia> Linstat Cheat Sheet (continued)

- > View all files in current folder:

```
ls
```

- > Go up one level in the filesystem (e.g. if we're in '/project_folder/code/model', we may want to go to 'project_folder/code'):

```
cd ..
```

- > More comprehensive list:

<https://gist.github.com/khazeamo/f762f532bfbc17d5bf396e9d4c2a9586>

- > Note: alternatively, you can access and navigate your Linux directory through Winstat (it is located in your Z drive)

julia> Linstat Cheat Sheet (continued)

> See jobs running on current Linstat server:

top

> To close top, press “q”

> Stop a job PID that is currently running:

kill PID

- > To run a script file on Slurm with N cores and M GB of memory:

```
ssubmit --cores=N --mem=Mg "julia_script.jl"
```

- > To figure out how much memory you need:
 1. Run your job on Linstat and check its memory use (RES)
 2. Then, kill your job and send it to Slurm
- > Slurm will send an email when your job is done running (either due to success or because it failed)
- > You can submit to specific partitions of the server using the option “-part”

julia> Slurm Cheat Sheet

- > Note: Most pre-compiled Julia package images won't work directly on Slurm; this means you will have to include the “-pkgimage=no” command in order to get your code to run:

```
ssubmit --cores=C --mem=Mg "julia --pkgimage=no my_julia_script.jl"
```

- > To see jobs currently running on Slurm, use:

```
squeue -a
```

- > To cancel your job:

```
scancel JOBID
```

```
julia> Linstat vs. Slurm
```

```
> Linstat is good for:
```

1. Interactive coding
2. Checking that your code runs on SSCC servers
3. Seeing how much memory you are using
4. Sending jobs to Slurm

```
> Slurm is goo for big jobs that need lots of cores and run for a long time
```